

Forward Deployed Engineer (FDE)

Embed with customers and turn their hardest problems into shipped AI solutions

Company	MONA AI GmbH — applied AI agents for enterprise
Location	Saarbrücken HQ + on-site at customers (travel up to ~30%)
Team	Delivery / Solutions Engineering · reports to Head of Delivery
Type	Full-time, permanent

About the role

FDEs are senior engineers who deploy directly into customer environments, scope ambiguous problems, and build and integrate AI-agent solutions against real data and systems. You own delivery from discovery to production hand-off, and you're the technical face to the customer.

What you'll do

- Scope customer problems on-site; translate vague requirements into a concrete technical plan.
- Build, integrate and deploy agentic workflows against customer data, APIs and internal systems.
- Design retrieval / RAG pipelines and evaluate LLM output quality with real test sets.
- Harden integrations: auth, error handling, observability, and security/PII handling.
- Run production hand-off, documentation and enablement; feed learnings back to Product.

Must-have qualifications

- 4+ years software engineering with strong Python (and SQL); production systems experience.
- Built and shipped LLM/agent or data-integration systems against messy real-world data.
- Solid on APIs, cloud (AWS/GCP/Azure), containers, and CI/CD.
- Customer-facing maturity: can run a technical workshop and say 'no' diplomatically.
- Fluent English; German a strong plus for on-site work in DE.

Nice to have

- Experience with RAG, vector databases, and LLM evaluation/guardrails.
- Background in regulated/enterprise environments (security, GDPR, audit).
- Prior consulting, solutions-engineering or FDE-style role.

Tools & stack

Python · SQL · LLM & agent frameworks · vector DBs (pgvector/Pinecone/Weaviate) · RAG & eval tooling · AWS/GCP/Azure · Docker · REST/gRPC · observability (logs/traces).

What success looks like (first 6 months)

At least one customer taken from discovery to production; a reusable integration pattern contributed back; measurable quality bar on agent outputs (eval pass-rate, not vibes).

For the agent: Dig into a project where requirements were ambiguous: how did they scope, what did they de-scope, how did they *evaluate* the AI output, and how did they handle security/PII? Red flags: only toy/demo projects, no eval methodology, or can't describe a real customer trade-off.